

The AI Engineer Profile Guide

A short review of the most common fixes I make to AI/ML Engineers' LinkedIn, resume and GitHub, with real before/after.



Jawad A.
Lead AI Platform Engineer · Toronto

As a Hiring Manager I've interviewed 100s of AI/ML Engineers and reviewed thousands of profiles over almost a decade.

CONTENTS

01	What I See Most the most common mistakes, counted	p. 02
02	Quick-Fix Checklist LinkedIn, resume and GitHub	p. 03
03	LinkedIn a full profile, then real headlines	p. 04
04	Resume real before/after	p. 06
05	GitHub three levels and a README	p. 08
06	Keep Going profiles worth studying, and where to start	p. 11

What I See Most

Counted line by line across the resumes, LinkedIn profiles and GitHub accounts of AI/ML job seekers who contacted me in September and October 2026.



Resume

N = 17 RESUMES

Claims without proof (a link, a number, a repo)		16/17 · 94%
No single target role		14/17 · 82%
An inconsistency a Hiring Manager would catch		14/17 · 82%
Best AI work buried below the fold		12/17 · 71%
Underselling in the first line		8/17 · 47%
Built evals, never claimed them		8/17 · 47%

LinkedIn

N = 20 PROFILES

No GitHub link in Contact info		16/20 · 80%
Current role blank or one line		12/20 · 60%
Headline isn't one AI/ML role		12/20 · 60%
No About section		6/20 · 30%
Top 3 skills unrelated to AI Engineering		5/20 · 25%

GitHub

N = 13 PROFILES · 63 REPOS

Showcased repos with no CI		58/63 · 92%
Showcased repos with no tests		55/63 · 87%
Active fewer than 26 of 52 weeks		10/13 · 77%
No profile README		9/13 · 69%

The pattern: these are mostly strong candidates. The problem is rarely the experience. It's that the evidence is missing, buried or contradicted, so a Hiring Manager skimming for 5 seconds never sees it.

The Quick-Fix Checklist

Tick each one that's true of your profile today.

LinkedIn

- | | |
|---|--|
| ☐ L1 Your headline names one role, plus your edge
"AI Engineer · LLM evals & cloud infra", not five keywords separated by pipes. | ☐ L2 The headline isn't a job request
"Seeking new grad roles" makes the job hunt your headline. The Open to Work frame is fine. |
| ☐ L3 Featured shows something you built
Something a Hiring Manager can click: a live demo, a shipped app, real users. | ☐ L4 Every role has a description, especially the current one
Recruiters search that text. 2–4 bullets with numbers, same as your resume. |
| ☐ L5 Your name and titles match your resume
A different surname or a stale role makes a recruiter wonder if it's you. | ☐ L6 About opens with what you build
Not "passionate recent graduate". Three lines: what you build, the proof, what you want next. |
| ☐ L7 Your top 3 skills are AI Engineering skills
They act as a second headline. "Figma, Product Strategy" pulls you away. | ☐ L8 No fancy Unicode fonts
Look-alike "bold" characters get stripped in exports and search, and break screen readers. |

Resume

- | | |
|---|--|
| ☐ R1 The title names the role you want
So does the file name.
"Resume_SWE_New_Grad_V6.pdf" gets read too. | ☐ R2 Your strongest AI work is in the top third
Not bullet 3 of job 2. |
| ☐ R3 Every skill appears in a bullet
If "MLOps" lives only in the skills list, it reads as keyword stuffing. | ☐ R4 Dates and titles survive a reference check
No unexplained overlaps. Keep official titles; add a descriptor. |
| ☐ R5 Strong metrics say how they were validated
0.90 AUC on 86 samples needs "held-out subjects, 5-fold CV". | ☐ R6 It reads cleanly as plain text
Paste it into a text editor. If it's scrambled, so is the ATS copy. |
| ☐ R7 Your title is backed by your recent roles
If your last 3 years don't show the AI work, I wonder where you learned it. Show the projects; link GitHub. | ☐ R8 Send a PDF, never a Word file
A .docx is editable, renders differently everywhere and looks unfinished. PDF is the standard. |

GitHub

- | | |
|---|--|
| ☐ G1 A one-line bio and a profile README
What you build, in one sentence. It's the first thing I read. | ☐ G2 2–4 pinned projects you're proud of
Not assignments, forks or dotfiles. Each states what it does and the result. |
| ☐ G3 READMEs lead with the result
The question, the result and a demo. Then how to run it. | ☐ G4 Activity shows steady building
Most weeks, all year. A burst right before the job hunt reads as exactly that. |

A Full Profile, Before and After

A new grad with LLM agent, eval and production cloud experience.
None of it showed.

FIG 3.1 - LEILA M. · NEW GRAD → AI ENGINEER

removed **added** **MOVED** **1** note, with an arrow to the lines it covers

I build and evaluate LLM agents

LM
#OPENTOWORK

Leila M.

Ex SWE @ Big Bank | ML & LLM Researcher | Seeking New Grad Roles

AI Engineer · LLM agents & evals · AWS / Terraform in production

Toronto, Ontario, Canada · 500+ connections

About

Recent Computer Science graduate passionate about AI/ML, cloud and data engineering. Seeking new grad or junior opportunities where I can learn and grow...

I build LLM agents and the evals that tell you whether they work.

Built an agentic recruitment assistant (LangChain, FastAPI) that 300+ job seekers have signed up for, with an eval pipeline that scores every release. Ran AWS infrastructure in production. Looking for an AI Engineer role on a team shipping LLM products.

Featured

Nothing featured

DEMO
Recruiting agent
2-min demo + repo

LIVE APP
Try it yourself
Paste a job post, get a tailored resume

SHIPPED
App Store
300+ sign-ups, 4.7★

Experience

U **Machine Learning Research Assistant** **Machine Learning Research Assistant**

University · May 2025 – Aug 2026 · **Part-time**

Contributed to LLM-based recruitment systems focused on fairness...

Built an agentic recruitment assistant and an eval pipeline over 100+ documents...

1 **Headline: one role plus an edge.**
"Ex", "Researcher" and "Seeking New Grad Roles" said where she'd been and what she wanted, not what she can do.

2 **About opens with what she built.**
"Passionate recent graduate" could be anyone. The agent and its users are the proof.

3 **Featured shows things to try.**
A demo, a live app, user numbers. It's the most visual spot on the profile, and it was empty.

4 **Experience says what she did.**
Fixed capitalization, marked part-time (it overlapped her co-ops), and leads with the system.

Real Headlines, Rewritten

Anonymized. Each becomes a single role plus an edge: check L1 on the checklist.

Sam A. • LinkedIn headline

- AI & ML Engineer | Machine Learning & Deep Learning | Medical Imaging & Healthcare AI | Data Analyst | Explainable AI

+ ML Engineer • Medical imaging & explainable deep learning

Five identities, including "Data Analyst", is the same as having none. Pick the one the experience supports.

Nadia S. • LinkedIn headline

- Ex-Big 4 | Machine Learning Associate | MMAI '25

+ AI Engineer • RAG & document AI • Shipped containerized LLM APIs

"Ex-Big 4" leads with someone else's brand. The fix leads with what she has shipped.

Omar Y. • LinkedIn headline

- Agentic AI Security Engineer at Gov. Research Lab role ended last year

+ AI Engineer | MSc | Prev. Agentic AI Security Engineer @ Gov. Research Lab

A stale "current" role raises a question in the first second. Leading with the target role and marking the old one "Prev." answers it.

Real Resumes, Before and After

Real resumes from people who messaged me, anonymized. Each shows the changes I suggested to them, tracked, and why.

FIG 4.1 — NADIA S. · MMAI GRAD, 1 YEAR IN → AI ENGINEERING

removed added MOVED 1 note, with an arrow to the lines it covers

Nadia S.

Toronto, ON · nadia.s@email.com · LinkedIn · [GitHub](#) · Portfolio

RELEVANT EXPERIENCE

AI Engineer, Big 4 firm Jan 2026 – Jun 2026

- Performed AI/ML data analysis, validation and quality assurance using Python and SQL, evaluating AI-generated outputs against business requirements...
- Applied SDLC best practices through SIT, UAT, smoke and regression testing, achieving a 95%+ pass rate...

Machine Learning Associate, AI research institute May 2025 – Sep 2025

- Owned end-to-end development of an API-based form-field extraction and classification solution, containerized with Docker; cut manual mapping from ~5 days to 3...
- Integrated an LLM (LLaMA) with a RAG framework for the document processing pipeline, reducing mapping errors by 70%

Artificial Intelligence Consultant Intern, Big Bank Sep 2024 – Apr 2025

- Developed and optimized a PII detection system using ML and NLP, rolled out across multiple internal departments...

~~Business Analyst, Startup – Oct 2021 – May 2022 – translated business requirements into technical specifications...~~

PROJECTS

- Portfolio chatbot agent, built with Claude · live: [nadia.s.netlify.app](#)
- Personal AI trainer prototype · captures goals and preferences to personalize a plan · repo · [demo]
- Mapper-AI · Python · repo · [one line on what it does and the result]

EDUCATION

Master of Management in Artificial Intelligence, business school
BS in Software Engineering

~~Certifications: Azure AI Fundamentals – IBM Data Analyst – Data Analyst (NPower) – Azure Fundamentals – Scrum Foundation – Intro to Python (MTA)~~

1 Put GitHub up top, next to LinkedIn.

Then get the activity up: anyone in AI these days is building every week.

2 Make the dates true.

A Hiring Manager will ask whether you're still there, and the end date answers it. Save the explanation for the conversation.

3 Get rid of the business analyst role.

Make room for what you've built with AI.

4 Add personal projects.

I want to see what you've built with AI, and that you're curious enough to keep up with this AI wave.

5 Remove the certifications.

What I look for is what you've built, not what you've certified.

FIG 4.2 — OMAR Y. • MSC RESEARCHER → AI ENGINEER (HIS ACTUAL REVISION)

Omar Y.

Ontario • omar.y@email.com • LinkedIn • GitHub • Medium

SUMMARY

AI Engineer with industry experience shipping production computer vision models and building LLM agent systems. MSc Computer Science, graduating Nov 2026.

INDUSTRY EXPERIENCE EXPERIENCE

~~Agentic AI Security Engineer~~ **AI Security Engineer, Agentic AI** • Gov. Research Lab Oct 2024 – Mar 2025

- ~~Engineered an automated workflow enabling LLMs to propose and execute command sequences in lateral movement attack scenarios~~

Built an LLM agent pipeline that proposes and executes commands in automated lateral movement scenarios, with step-level logging

- Benchmarked 10 LLMs across 5 temperature settings, analyzing behavioral differences between configurations

~~Computer Vision Engineer~~ **AI Engineer, Computer Vision** • AI Startup Aug 2024 – Aug 2025

- Trained and shipped** production detection, classification and OCR models; improved a parking-spot detection system deployed city-wide
- Built a vector-search retrieval system over 6M images; **evaluated** VLMs, CLIP-style models and LLMs for production, trading off accuracy, latency and cost

~~SOC Analyst Intern, 2023–Telecommunications Internship, 2022~~

~~Academic Experience: 3 research assistant roles, 9 bullets~~

PROJECTS

- LLM Guardrail Evaluation Framework** • 6 guardrail models vs 4 attack families • GitHub
- Open-source project-matching platform** • Supabase, LLM API, CI/CD • Live • GitHub
- ~~Power analysis attacks • Wireshark forensics • NYC taxi data cleaning • University website backend • Query optimization~~

SKILLS

AI/ML: PyTorch, HuggingFace, LLM APIs, LLM agents, embeddings & vector search, model evaluation

~~Networks & Security: TCP/IP, QKD • Cybersecurity tools: QRadar, VirusTotal, Nmap~~

EDUCATION MOVED FROM THE TOP

MSc Computer Science (thesis: LLM guardrail security) 2026

1 Go the AI Engineer route.

Research needs at least a PhD, is very competitive, and has fewer roles in Canada. This resume is already more Engineer.

2 Use "AI" in the title.

The work already uses ML techniques; the title should say so.

3 Concise bullets in AI Engineer language.

Show the cloud and Kubernetes experience, PyTorch, and the SDLC with AI.

4 Cut the off-target roles.

The telecom and analyst internships, and the academic section, go.

5 Three projects you're proudest of.

Built with AI, each with web, GitHub or demo links.

6 Simplify.

"Industry Experience" becomes Experience; skills read like an AI Engineer's.

Weak, Okay, Strong

The same person's GitHub at three levels, one per page. Under each is what I conclude in the first 10 seconds.

- WEAK

The screenshot shows a GitHub profile for user **oy2001**. The profile is sparse and lacks a bio. The pinned repositories are:

- cs3307-assignment3**: No description, Java.
- transformers**: Forked from huggingface/transformers, Python.
- dotfiles**: No description, Shell.
- project-final**: No description, JavaScript.

The contribution graph shows 41 contributions in the last year, represented by a grid of green squares.

What I conclude: no evidence they build anything outside class.

What would move it to okay

- Write a one-line bio.**
What you build, in one sentence. It's the first line I read. (G1)
- Unpin the assignment, the fork and the dotfiles.**
Pin 2–4 things you built yourself, even small ones. (G2)
- Give every pinned repo a description.**
What it does and the result, in one line.
- Commit something most weeks.**
A small project you return to beats a burst before the job hunt. (G4)

~ OKAY

oy

oy2001
ML engineer · Toronto

README.md

No profile README

Pinned

guardrail-bench
Benchmarking guardrails
Python

plate-ocr
License plate OCR
Python

rag-playground
No description
Python

portfolio
Personal site
TypeScript

349 contributions in the last year

What I conclude: real projects and real activity, but the activity comes in bursts and mostly happened in the last month. It reads as "started building for the job hunt". It doesn't hurt, but it won't make you stand out.

What would move it to strong

- Add a profile README.**
Two or three lines: what you build, what you're working on now, one link to writing.
- Turn descriptions into results.**
"Benchmarking guardrails" becomes "11 guardrail models vs 6 attack families, results table, 1-command repro".
- Lead each README with the result and a demo.**
The question, the answer, a 2-minute video. Then how to run it. (G3)
- Make the activity steady, not seasonal.**
Most weeks, all year. It's the one thing with no shortcut. (G4)

+ STRONG

oy

oy2001

AI Engineer · LLM security & evals · Toronto

oy2001 / README.md

I build ways to test whether LLM systems are safe to ship.

Now: open-source red-team harness for tool-using agents

Latest write-up: [Your guardrail model has a blind spot](#)

Stack: PyTorch · vLLM · FastAPI · Docker · GitHub Actions

Pinned

guardrail-bench

11 guardrail models vs 6 attack families. Results table, write-up, 1-command repro.

Python 48

agent-redteam

Red-team harness for tool-using agents. ▶ 2-min demo.

Python 21

rag-eval-kit

Faithfulness and retrieval evals that run as CI checks.


Python 9

plate-ocr

License-plate OCR, from detection to recognition. Live demo.

Python 6

612 contributions in the last year



What I conclude: a builder with a specialty. I know what to ask about before the interview starts.

A Project README, Before and After

- BEFORE

```
# guardrail-bench

run: python main.py
```

1 Nothing to judge.

I can't tell what it does, whether it works, or why it matters, and I won't clone it to find out.

+ AFTER

```
# guardrail-bench
Which LLM guardrails actually stop attacks?
11 models × 6 attack families, one command.

Demo: 2-min video · Write-up · MSc thesis

## Results (attack success rate, lower is better)
| Guardrail | Jailbreak | Prompt inj. | Encoding |
|-----|-----|-----|-----|
| model-a | 4% | 31% | 58% |
| model-b | 9% | 12% | 22% |

## Run it
make bench MODEL=model-a

## How it works
attacks/ → harness/ → scorers/ → report.md
```

THE AI ENGINEER PROFILE GUIDE · GITHUB

10 / 11

Profiles Worth Studying

Exceptional, and not the bar for an interview. Look at the pattern: a clear one-liner, a few flagship projects, writing that explains them, steady output.

PROFILE	WHAT TO NOTICE
 simonw	Ships small tools constantly (llm, datasette), and every one has a blog post explaining it.
 karpthy	A few pinned repos (nanoGPT), each one minimal, readable and focused on teaching.
 rasbt	A profile README that routes you to the work. LLMs-from-scratch is a model flagship project.
 eugeneyan	Pairs repos with writing on applied ML. Proof that writing about your work counts.
 hame1smu	Practical writing on LLM evals, the most under-claimed skill on resumes.

How to Use This

1. **Start with the checklist on page 3.** Every empty box is a fix you can make this week.
2. **Fix your headline and resume title first.** They're the first thing a Hiring Manager reads, and they set how everything after them is read.
3. **Then add proof.** A link, a number, a demo or a repo for every claim that matters.
4. **Ask someone in the role you want to read it for 5 seconds,** then tell you what you do. If their answer isn't your target role, keep going.

Questions or feedback on this guide? Message me on [LinkedIn](#).