

The AI Engineer Profile Playbook

What a Hiring Manager checks in your LinkedIn, resume and GitHub, and how to fix each one, with real before/after.



Jawad A.
Lead AI Platform Engineer, Toronto. As a Hiring Manager I've interviewed 100s of AI/ML Engineers over almost a decade.

CONTENTS

01	What I see most the most common mistakes, counted	p. 02
02	18 quick fixes LinkedIn, resume and GitHub	p. 03
03	LinkedIn a full profile, then five headlines	p. 04
04	Resume three before/after	p. 06
05	GitHub three levels and a README	p. 09
06	Next step profiles to study, and the full audit	p. 11

HOW TO READ THE FIGURES

removed

added

MOVED


1 note, with an arrow to the lines it covers

What I see most

Counted line by line across the resumes, LinkedIn profiles and GitHub accounts of AI/ML job seekers who contacted me in September and October 2026.



Resume

N = 17 RESUMES

Claims without proof (a link, a number, a repo)		16/17 · 94%
No single target role		14/17 · 82%
An inconsistency a Hiring Manager would catch		14/17 · 82%
Best AI work buried below the fold		12/17 · 71%
Underselling in the first line		8/17 · 47%
Built evals, never claimed them		8/17 · 47%

LinkedIn

N = 20 PROFILES

No GitHub link in Contact info		16/20 · 80%
Current role blank or one line		12/20 · 60%
Headline isn't one AI/ML role		12/20 · 60%
No About section		6/20 · 30%
Top 3 skills unrelated to AI Engineering		5/20 · 25%

GitHub

N = 13 PROFILES · 63 REPOS

Showcased repos with no CI		58/63 · 92%
Showcased repos with no tests		55/63 · 87%
Active fewer than 26 of 52 weeks		10/13 · 77%
No profile README		9/13 · 69%

The pattern: these are mostly strong candidates. The problem is rarely the experience. It's that the evidence is missing, buried or contradicted, so a Hiring Manager skimming for 5 seconds never sees it.

18 quick fixes

Tick each one that's true of your profile today. Every empty box is a fix you can make this week.

LinkedIn

- | | |
|---|--|
| ☐ L1 Your headline names one role, plus your edge
"AI Engineer · LLM evals & cloud infra", not five keywords separated by pipes. | ☐ L2 The headline isn't a job request
"Seeking new grad roles" makes the job hunt your headline. The Open to Work frame is fine. |
| ☐ L3 Featured shows something you built
Something a Hiring Manager can click: a live demo, a shipped app, real users. | ☐ L4 Every role has a description, especially the current one
Recruiters search that text. 2–4 bullets with numbers, same as your resume. |
| ☐ L5 Your name and titles match your resume
A different surname or a stale role makes a recruiter wonder if it's you. | ☐ L6 About opens with what you build
Not "passionate recent graduate". Three lines: what you build, the proof, what you want next. |
| ☐ L7 Your top 3 skills are AI Engineering skills
They act as a second headline. "Figma, Product Strategy" pulls you away. | ☐ L8 No fancy Unicode fonts
Look-alike "bold" characters get stripped in exports and search, and break screen readers. |

Resume

- | | |
|---|---|
| ☐ R1 The title names the role you want
So does the file name.
"Resume_SWE_New_Grad_V6.pdf" gets read too. | ☐ R2 Your strongest AI work is in the top third
Not bullet 3 of job 2. |
| ☐ R3 Every skill appears in a bullet
If "MLOps" lives only in the skills list, it reads as keyword stuffing. | ☐ R4 Dates and titles survive a reference check
No unexplained overlaps. Keep official titles; add a descriptor. |
| ☐ R5 Strong metrics say how they were validated
0.90 AUC on 86 samples needs "held-out subjects, 5-fold CV". | ☐ R6 It reads cleanly as plain text
Paste it into a text editor. If it's scrambled, so is the ATS copy. |

GitHub

- | | |
|---|--|
| ☐ G1 A one-line bio and a profile README
What you build, in one sentence. It's the first thing I read. | ☐ G2 2–4 pinned projects you're proud of
Not assignments, forks or dotfiles. Each states what it does and the result. |
| ☐ G3 READMEs lead with the result
The question, the result and a demo. Then how to run it. | ☐ G4 Activity shows steady building
Most weeks, all year. A burst right before the job hunt reads as exactly that. |

A full profile, before and after

A new grad with LLM agent, eval and production cloud experience.
None of it showed.

FIG 3.1 - LEILA M. · NEW GRAD → AI ENGINEER

~~removed~~ **added** **MOVED** **1** note, with an arrow to the lines it covers

I build and evaluate LLM agents

LM
#OPENTOWORK

Leila M.

~~Ex SWE @ Big Bank | ML & LLM Researcher | Seeking New Grad Roles~~

AI Engineer · LLM agents & evals · AWS / Terraform in production

Toronto, Ontario, Canada · 500+ connections

About

~~Recent Computer Science graduate passionate about AI/ML, cloud and data engineering. Seeking new grad or junior opportunities where I can learn and grow...~~

I build LLM agents and the evals that tell you whether they work.

Built an agentic recruitment assistant (LangChain, FastAPI) that 300+ job seekers have signed up for, with an eval pipeline that scores every release. Ran AWS infrastructure in production. Looking for an AI Engineer role on a team shipping LLM products.

Featured

~~Nothing featured~~

DEMO Recruiting agent	LIVE APP Try it yourself	SHIPPED App Store
2-min demo + repo	Paste a job post, get a tailored resume	300+ sign-ups, 4.7★

Experience

U ~~Machine learning Research Assistant~~ **Machine Learning Research Assistant**

University · May 2025 – Aug 2026 · **Part-time**

~~Contributed to LLM-based recruitment systems focused on fairness...~~

Built an agentic recruitment assistant and an eval pipeline over 100+ documents...

1 **Headline: one role plus an edge.**
"Ex", "Researcher" and "Seeking New Grad Roles" said where she'd been and what she wanted, not what she can do.

2 **About opens with what she built.**
"Passionate recent graduate" could be anyone. The agent and its users are the proof.

3 **Featured shows things to try.**
A demo, a live app, user numbers. It's the most visual spot on the profile, and it was empty.

4 **Experience says what she did.**
Fixed capitalization, marked part-time (it overlapped her co-ops), and leads with the system.

Five real headlines, rewritten

Anonymized. Each becomes a single role plus an edge: check L1 on the checklist.

Sam A. • LinkedIn headline

- AI & ML Engineer | Machine Learning & Deep Learning | Medical Imaging & Healthcare AI | Data Analyst | Explainable AI

+ ML Engineer • Medical imaging & explainable deep learning

Five identities, including "Data Analyst", is the same as having none. Pick the one the experience supports.

Marcus K. • LinkedIn headline

- AI & ML @ Consultancy | M.S. Computer Science

+ AI Engineering Lead • Production RAG & Text-to-SQL • Led teams of 6

The original names an employer and a degree. The fix names the role, and the proof that answers "can he lead engineers?"

Nadia S. • LinkedIn headline

- Ex-Big 4 | Machine Learning Associate | MMAI '25

+ AI Engineer • RAG & document AI • Shipped containerized LLM APIs

"Ex-Big 4" leads with someone else's brand. The fix leads with what she has shipped.

Omar Y. • LinkedIn headline

- Agentic AI Security Engineer at Gov. Research Lab role ended last year

+ AI Engineer • LLM security & evals • MSc thesis on guardrails

A stale "current" role raises a question in the first second. The fix turns it into a specialty.

Rami A. • LinkedIn headline

- Senior Machine Learning Engineer @ Big Bank accurate, forgettable

+ Senior ML Engineer • LLM workflows & ML observability • Shipped a paid iOS app

The fix adds the two things that set him apart, including an app with paying users.

Three resumes, before and after

Each is one document with the changes tracked: what was removed, what was added, and why.

FIG 4.1 - MAYA R. · NEW GRAD → ML INFRASTRUCTURE

removed **added** **MOVED** **1** note, with an arrow to the lines it covers

Maya R.

Software Engineer · New Grad

ML Infrastructure Engineer · LLM observability · Python, Go

Toronto, ON · maya.r@email.com · [github.com/mayar](#) · [mayar.dev](#)

File: ~~Resume_SWE_New_Grad_V6.pdf~~ [Maya_R_ML_Infrastructure.pdf](#)

EXPERIENCE

Social Media Co. · Backend Engineering Intern, Ads ML Platform Jun-Sep 2025

- Built a unified LLM observability library (latency, tokens, cost) across 3 model providers, adopted by 5 ML services for on-call debugging and cost tracking
- ~~Developed a Unified LLM Observability Tool, standardizing metric collection... (was bullet 3)~~
- Designed the platform's first integration test framework (gRPC, permissioned access to real services), with Slack alerting on regressions
- Replaced Graphite with a Prometheus/OpenTelemetry metrics library for high-throughput ML services

Fintech Co. · Software Engineer II ~~(role eliminated)~~ Jun-Sep 2026 SWE Co-op Sep-Dec 2025

- Added Sentry and Datadog observability to previously unmonitored services, including an MCP server, and prototyped AI-assisted incident triage
- Cut table load time from 6s to under 1s with a decorator-based caching layer

Big Bank · Software Developer Intern (2 terms) 2024

SELECTED WORK

- Multi-agent RL for cooperation, co-author, top-tier ML conference 2026 · [paper](#) · [code](#)
- tinyglot: character-level Transformer, 10 languages · [live demo](#) · [write-up](#)

EDUCATION

BCS (Hons.), Co-op MOVED FROM THE TOP 2026

- ~~Coursework: Distributed Systems, Networks, NLP, Neural Networks, RL~~

1

Title yourself as the role you want.

Recruiters read the title and the file name before anything else. "New grad" says stage, not skill.

2

Lead with the strongest AI evidence.

The LLM observability work was bullet 3 of job 2. For an ML infra role, it's the headline, with numbers.

3

Don't explain a gap on paper.

If a Hiring Manager wants to know, they'll ask. Have a one-line answer ready.

4

Link everything.

A paper or project nobody can click on counts for little.

5

Cut the coursework line.

It takes space that shipped work should have. Education moves to the bottom.

THE AI ENGINEER PROFILE PLAYBOOK · RESUME

06 / 11

FIG 4.2 – MARCUS K. • DATA SCIENTIST → AI ENGINEERING LEAD

Marcus K.

Data Scientist

AI Engineering Lead · Production RAG & LLM systems · 8 yrs

Toronto, ON · marcus.k@email.com · github.com/marcusk

SUMMARY

Data Scientist and tech lead with 8 years of experience designing scalable AI & ML solutions...

AI Engineering lead who ships LLM systems to production for enterprise clients and leads engineers. Has led teams of up to 6 and turned research-grade prototypes into evaluated, monitored services.

EXPERIENCE

Consultancy · AI Consultant (AI Engineering lead)2024 - now

Built a production RAG architecture... developing ETL and chunking pipelines, embedding data into a vector store...

Led design and delivery of a production RAG system for a pharma client (Databricks, vector search); built a retrieval eval harness (recall@k, faithfulness) used to gate every release

Shipped Text-to-SQL for actuaries on Azure, saving several hours per analyst each week; ran stakeholder workshops to scope it

E-commerce Co. · Data Scientist2022 - 2024

Consolidated four legacy LTV systems into one modular framework: \$5M/yr savings, guiding \$600M in spend

Consultancy · Data Scientist (Tech Lead)2018 - 2022

Led 6 engineers and BAs to ship an order distribution optimizer: \$3M/yr freight savingsALSO IN SUMMARY

SKILLS

LLM systems: RAG, Text-to-SQL, LangGraph, evals and tracing

Engineering: Python, FastAPI, Docker, CI/CD, Spark, Airflow · Databricks, GCP, AWS, Azure

Tableau, Looker, Plotly, Seaborn, D3.js, Hive

Relevant Courses: ML, NLP, Bayesian Statistics...

1

The title is a filter.

To a recruiter, "Data Scientist" reads as "can't lead engineers or work with PMs". Name the role you're applying for.

2

Answer the biggest doubt first.

The proof of leading engineers was buried in 2018. It goes in the summary.

3

Show the engineering, not just the outcome.

Say how quality was measured. An eval harness is the most under-claimed skill there is. Only claim it if you built it.

4

Lead with the LLM stack.

BI-heavy tools pull the reader back toward "analyst".

5

No "Relevant Courses" at 8 years.

It undercuts the seniority.

FIG 4.3 — OMAR Y. • MSC RESEARCHER → AI ENGINEER

Omar Y. Hassan

Was "Omar Y." here and "Omar H." on LinkedIn

~~MSc CS candidate — looking for research or engineering roles~~

AI Engineer • LLM security & evals • Eligible to work in Canada

Toronto, ON • omar.h@email.com • github.com/oy2001

EXPERIENCE

Gov. Research Lab • Agentic AI Security Engineer *Oct 2024 – Mar 2025 • Remote*

- Built an agentic harness where LLMs plan and execute lateral-movement attacks; benchmarked 10 models × 5 temperatures across 3 autonomy modes

AI Startup • Computer Vision Engineer (Deep Learning, VLMs) *Aug 2024 – Aug 2025 • Cairo (on-site)*

- Shipped production detection and OCR models; built CLIP-style retrieval over 6M images

~~SOC Analyst Intern — Telecom Internship, 2022–23~~

~~Academic Experience: 4 research assistant roles...~~

PROJECTS

- guardrail-bench: 11 LLM guardrail models vs 6 attack families • repo • write-up • thesis
- agent-redteam: open-source harness for testing tool-using agents • demo video
- rag-eval-kit: faithfulness and retrieval evals as CI checks • repo
- ~~Wireshark forensics — NYC taxi cleaning — University website backend — Fantasy football analysis~~

EDUCATION

MSc Computer Science (thesis: LLM guardrail security) *2026*

1 One name everywhere.
A different name on the resume than on LinkedIn means a recruiter can't connect the two.

2 Pick one: research or engineering.
Research roles mostly need a PhD and are scarce in Canada. Pick the role the evidence supports, then specialize.

3 Explain overlapping jobs.
Two countries at once looks like an error without "Remote". Keep official titles and add a descriptor, so it survives a reference check.

4 Cut off-target roles.
It makes room for the specialty.

5 Make the thesis the projects.
The AI security and evals work was one line under Education. Now it's three linked projects.

THE AI ENGINEER PROFILE PLAYBOOK • RESUME

08 / 11

Weak, okay, strong

The same person's GitHub at three levels. Under each is what I conclude in the first 10 seconds.

- WEAK

oy2001
(no bio)

README.md
No profile README

Pinned

- cs3307-assignment3**
No description
Java
- transformers**
Forked from huggingface/transformers
Python
- dotfiles**
No description
Shell
- project-final**
No description
JavaScript

41 contributions in the last year

A grid showing 41 contributions in the last year, with most activity concentrated in the last month.

What I conclude: no evidence they build anything outside class.

~ OKAY

oy2001
ML engineer · Toronto

README.md
No profile README

Pinned

- guardrail-bench**
Benchmarking guardrails
Python
- plate-ocr**
License plate OCR
Python
- rag-playground**
No description
Python
- portfolio**
Personal site
TypeScript

349 contributions in the last year

A grid showing 349 contributions in the last year, with activity spread across the year but with a notable burst in the last month.

What I conclude: real projects and real activity, but the activity comes in bursts and mostly happened in the last month. It reads as "started building for the job hunt". It doesn't hurt, but it won't make you stand out.

+ STRONG

oy

oy2001

AI Engineer · LLM security & evals · Toronto

oy2001 / README.md

I build ways to test whether LLM systems are safe to ship.

Now: open-source red-team harness for tool-using agents

Latest write-up: [Your guardrail model has a blind spot](#)

Stack: PyTorch · vLLM · FastAPI · Docker · GitHub Actions

Pinned

guardrail-bench

11 guardrail models vs 6 attack families. Results table, write-up, 1-command repro.

Python ★ 48

agent-redteam

Red-team harness for tool-using agents. ▶ 2-min demo.

Python ★ 21

rag-eval-kit

Faithfulness and retrieval evals that run as CI checks.

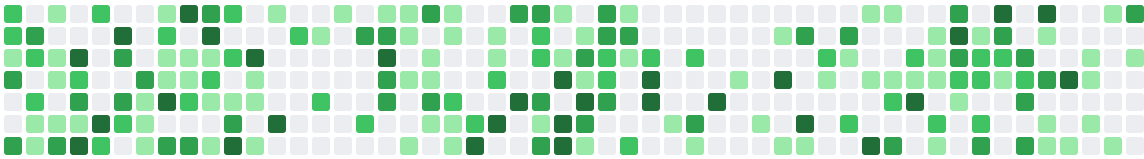
Python ★ 9

plate-ocr

License-plate OCR, from detection to recognition. Live demo.

Python ★ 6

612 contributions in the last year



What I conclude: a builder with a specialty. I know what to ask about before the interview starts.

A project README, before and after

- BEFORE

```
# guardrail-bench

run: python main.py
```

1 Nothing to judge.

I can't tell what it does, whether it works, or why it matters, and I won't clone it to find out.

+ AFTER

```
# guardrail-bench
Which LLM guardrails actually stop attacks?
11 models × 6 attack families, one command.

Demo: 2-min video · Write-up · MSc thesis

## Results (attack success rate, lower is better)
| Guardrail | Jailbreak | Prompt inj. | Encoding |
|-----|-----|-----|-----|
| model-a | 4% | 31% | 58% |
| model-b | 9% | 12% | 22% |

## Run it
make bench MODEL=model-a

## How it works
attacks/ → harness/ → scorers/ → report.md
```

THE AI ENGINEER PROFILE PLAYBOOK · GITHUB

10 / 11

Profiles worth studying

Exceptional, and not the bar for an interview. Look at the pattern: a clear one-liner, a few flagship projects, writing that explains them, steady output.

PROFILE	WHAT TO NOTICE
github.com/simonw	Ships small tools constantly (llm, datasette), and every one has a blog post explaining it.
github.com/karpathy	A few pinned repos (nanoGPT), each one minimal, readable and focused on teaching.
github.com/rasbt	A profile README that routes you to the work. LLMs-from-scratch is a model flagship project.
github.com/eugeneyan	Pairs repos with writing on applied ML. Proof that writing about your work counts.
github.com/hamelsmu	Practical writing on LLM evals, the most under-claimed skill on resumes.

Want these checks run on your profile?

This playbook is the sample. An audit applies every check to **your** LinkedIn, resume and GitHub, ranks the fixes by impact and picks one target role your evidence supports. I give the feedback; you make the changes.

TIER	INCLUDES	CAD
Audit WRITTEN · 72 H RECOMMENDED	Line-by-line audit of LinkedIn, resume and GitHub. Every fix ranked by impact. One target role and a top-3 project plan. 1 review of your revision.	\$249
Rebuild LIVE · ~2 WEEKS	Everything in the Audit, plus 2 × 45-min walkthroughs, GitHub pins and README guidance, a mock interview, and 2 reviews within 30 days.	\$749
Lead Rebuild LIVE · SENIOR ROLES	Everything in the Rebuild, plus a leadership story (leading engineers, working with PMs), system-design and behavioral mocks, and 2 reviews within 60 days.	\$1,250

Audit guarantee: fewer than 10 specific, actionable fixes and you get your money back. Book at jateeq.github.io/ai-engineer-audit/ or message me on [LinkedIn](#).

Data: 17 resumes, 20 LinkedIn profiles and 13 GitHub accounts (63 showcased repos) from AI/ML job seekers, Sep-Oct 2026. GitHub figures come from public data via the GitHub API.

Every example is anonymized: names, employers, schools and numbers changed. Contribution graphs are illustrative.

Playbook v1.0 · October 2026 · Set in Inter, Source Serif 4 and JetBrains Mono.