

# The AI Engineer Profile Playbook

What a Hiring Manager checks in your LinkedIn, resume and GitHub, and how to fix each one, with real before/after.



**Jawad A.**  
Lead AI Platform Engineer, Toronto. As a Hiring Manager I've interviewed 100s of AI/ML Engineers over almost a decade.

CONTENTS

|    |                                                          |       |
|----|----------------------------------------------------------|-------|
| 01 | <b>What I see most</b> the most common mistakes, counted | p. 02 |
| 02 | <b>18 quick fixes</b> LinkedIn, resume and GitHub        | p. 03 |
| 03 | <b>LinkedIn</b> a full profile, then five headlines      | p. 04 |
| 04 | <b>Resume</b> three before/after                         | p. 06 |
| 05 | <b>GitHub</b> three levels and a README                  | p. 09 |
| 06 | <b>Next step</b> profiles to study, and the full audit   | p. 11 |

HOW TO READ THE FIGURES

removed

added

MOVED


note, with an arrow to the lines it covers

# What I see most

Counted line by line across the resumes, LinkedIn profiles and GitHub accounts of AI/ML job seekers who contacted me in September and October 2026.



## Resume

N = 17 RESUMES

|                                                 |                                                                                    |             |
|-------------------------------------------------|------------------------------------------------------------------------------------|-------------|
| Claims without proof (a link, a number, a repo) |  | 16/17 · 94% |
| No single target role                           |  | 14/17 · 82% |
| An inconsistency a Hiring Manager would catch   |  | 14/17 · 82% |
| Best AI work buried below the fold              |  | 12/17 · 71% |
| Underselling in the first line                  |  | 8/17 · 47%  |
| Built evals, never claimed them                 |  | 8/17 · 47%  |

## LinkedIn

N = 20 PROFILES

|                                          |                                                                                      |             |
|------------------------------------------|--------------------------------------------------------------------------------------|-------------|
| No GitHub link in Contact info           |  | 16/20 · 80% |
| Current role blank or one line           |  | 12/20 · 60% |
| Headline isn't one AI/ML role            |  | 12/20 · 60% |
| No About section                         |  | 6/20 · 30%  |
| Top 3 skills unrelated to AI Engineering |  | 5/20 · 25%  |

## GitHub

N = 13 PROFILES · 63 REPOS

|                                  |                                                                                      |             |
|----------------------------------|--------------------------------------------------------------------------------------|-------------|
| Showcased repos with no CI       |  | 58/63 · 92% |
| Showcased repos with no tests    |  | 55/63 · 87% |
| Active fewer than 26 of 52 weeks |  | 10/13 · 77% |
| No profile README                |  | 9/13 · 69%  |

**The pattern:** these are mostly strong candidates. The problem is rarely the experience. It's that the evidence is missing, buried or contradicted, so a Hiring Manager skimming for 5 seconds never sees it.

# 18 quick fixes

Tick each one that's true of your profile today. Every empty box is a fix you can make this week.

## LinkedIn

- |                                                                                                                                                                   |                                                                                                                                                                            |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <input type="checkbox"/> L1 <b>Your headline names one role, plus your edge</b><br>"AI Engineer · LLM evals & cloud infra", not five keywords separated by pipes. | <input type="checkbox"/> L2 <b>The headline isn't a job request</b><br>"Seeking new grad roles" makes the job hunt your headline. The Open to Work frame is fine.          |
| <input type="checkbox"/> L3 <b>Featured shows something you built</b><br>Something a Hiring Manager can click: a live demo, a shipped app, real users.            | <input type="checkbox"/> L4 <b>Every role has a description, especially the current one</b><br>Recruiters search that text. 2–4 bullets with numbers, same as your resume. |
| <input type="checkbox"/> L5 <b>Your name and titles match your resume</b><br>A different surname or a stale role makes a recruiter wonder if it's you.            | <input type="checkbox"/> L6 <b>About opens with what you build</b><br>Not "passionate recent graduate". Three lines: what you build, the proof, what you want next.        |
| <input type="checkbox"/> L7 <b>Your top 3 skills are AI Engineering skills</b><br>They act as a second headline. "Figma, Product Strategy" pulls you away.        | <input type="checkbox"/> L8 <b>No fancy Unicode fonts</b><br>Look-alike "bold" characters get stripped in exports and search, and break screen readers.                    |

## Resume

- |                                                                                                                                               |                                                                                                                                                   |
|-----------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|
| <input type="checkbox"/> R1 <b>The title names the role you want</b><br>So does the file name.<br>"Resume_SWE_New_Grad_V6.pdf" gets read too. | <input type="checkbox"/> R2 <b>Your strongest AI work is in the top third</b><br>Not bullet 3 of job 2.                                           |
| <input type="checkbox"/> R3 <b>Every skill appears in a bullet</b><br>If "MLOps" lives only in the skills list, it reads as keyword stuffing. | <input type="checkbox"/> R4 <b>Dates and titles survive a reference check</b><br>No unexplained overlaps. Keep official titles; add a descriptor. |
| <input type="checkbox"/> R5 <b>Strong metrics say how they were validated</b><br>0.90 AUC on 86 samples needs "held-out subjects, 5-fold CV". | <input type="checkbox"/> R6 <b>It reads cleanly as plain text</b><br>Paste it into a text editor. If it's scrambled, so is the ATS copy.          |

## GitHub

- |                                                                                                                                         |                                                                                                                                                        |
|-----------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------|
| <input type="checkbox"/> G1 <b>A one-line bio and a profile README</b><br>What you build, in one sentence. It's the first thing I read. | <input type="checkbox"/> G2 <b>2–4 pinned projects you're proud of</b><br>Not assignments, forks or dotfiles. Each states what it does and the result. |
| <input type="checkbox"/> G3 <b>READMEs lead with the result</b><br>The question, the result and a demo. Then how to run it.             | <input type="checkbox"/> G4 <b>Activity shows steady building</b><br>Most weeks, all year. A burst right before the job hunt reads as exactly that.    |

# A full profile, before and after

A new grad with LLM agent, eval and production cloud experience.  
None of it showed.

FIG 3.1 - LEILA M. · NEW GRAD → AI ENGINEER

~~removed~~ **added** **MOVED**  note, with an arrow to the lines it covers



**I build and evaluate LLM agents**

**Leila M.**

~~Ex SWE @ Big Bank | ML & LLM Researcher | Seeking New Grad Roles~~

AI Engineer · LLM agents & evals · AWS / Terraform in production

Toronto, Ontario, Canada · 500+ connections

**About**

~~Recent Computer Science graduate passionate about AI/ML, cloud and data engineering. Seeking new grad or junior opportunities where I can learn and grow...~~

I build LLM agents and the evals that tell you whether they work.

Built an agentic recruitment assistant (LangChain, FastAPI) that 300+ job seekers have signed up for, with an eval pipeline that scores every release. Ran AWS infrastructure in production. Looking for an AI Engineer role on a team shipping LLM products.

**Featured**

~~Nothing featured~~

**DEMO**  
Recruiting agent

2-min demo + repo

**LIVE APP**  
Try it yourself

Paste a job post, get a tailored resume

**SHIPPED**  
App Store

300+ sign-ups, 4.7★

**Experience**

**U** ~~Machine learning Research Assistant~~ **Machine Learning Research Assistant**

University · May 2025 – Aug 2026 · **Part-time**

~~Contributed to LLM-based recruitment systems focused on fairness...~~

Built an agentic recruitment assistant and an eval pipeline over 100+ documents...

**1** **Headline: one role plus an edge.**  
"Ex", "Researcher" and "Seeking New Grad Roles" said where she'd been and what she wanted, not what she can do.

**2** **About opens with what she built.**  
"Passionate recent graduate" could be anyone. The agent and its users are the proof.

**3** **Featured shows things to try.**  
A demo, a live app, user numbers. It's the most visual spot on the profile, and it was empty.

**4** **Experience says what she did.**  
Fixed capitalization, marked part-time (it overlapped her co-ops), and leads with the system.

## Five real headlines, rewritten

Anonymized. Each becomes a single role plus an edge: check L1 on the checklist.

Sam A. • LinkedIn headline

- AI & ML Engineer | Machine Learning & Deep Learning | Medical Imaging & Healthcare AI | Data Analyst | Explainable AI

+ ML Engineer • Medical imaging & explainable deep learning

# Five identities, including "Data Analyst", is the same as having none. Pick the one the experience supports.

Marcus K. • LinkedIn headline

- AI & ML @ Consultancy | M.S. Computer Science

+ AI Engineering Lead • Production RAG & Text-to-SQL • Led teams of 6

# The original names an employer and a degree. The fix names the role, and the proof that answers "can he lead engineers?"

Nadia S. • LinkedIn headline

- Ex-Big 4 | Machine Learning Associate | MMAI '25

+ AI Engineer • RAG & document AI • Shipped containerized LLM APIs

# "Ex-Big 4" leads with someone else's brand. The fix leads with what she has shipped.

Omar Y. • LinkedIn headline

- Agentic AI Security Engineer at Gov. Research Lab role ended last year

+ AI Engineer • LLM security & evals • MSc thesis on guardrails

# A stale "current" role raises a question in the first second. The fix turns it into a specialty.

Rami A. • LinkedIn headline

- Senior Machine Learning Engineer @ Big Bank accurate, forgettable

+ Senior ML Engineer • LLM workflows & ML observability • Shipped a paid iOS app

# The fix adds the two things that set him apart, including an app with paying users.

# Three resumes, before and after

Each is one document with the changes tracked: what was removed, what was added, and why.

FIG 4.1 - MAYA R. · NEW GRAD → ML INFRASTRUCTURE

**removed**   **added**   **MOVED**   **1** note, with an arrow to the lines it covers

## Maya R.

~~Software Engineer · New Grad~~

ML Infrastructure Engineer · LLM observability · Python, Go

Toronto, ON · maya.r@email.com · [github.com/mayar](#) · [mayar.dev](#)

File: ~~Resume\_SWE\_New\_Grad\_V6.pdf~~ [Maya\\_R\\_ML\\_Infrastructure.pdf](#)

EXPERIENCE

Social Media Co. · Backend Engineering Intern, Ads    *Jun-Sep 2025*  
ML Platform

Built a unified LLM observability library (latency, tokens, cost) across 3 model providers, adopted by 5 ML services for on-call debugging and cost tracking

~~Developed a Unified LLM Observability Tool, standardizing metric collection... (was bullet 3)~~

Designed the platform's first integration test framework (gRPC, permissioned access to real services), with Slack alerting on regressions

Replaced Graphite with a Prometheus/OpenTelemetry metrics library for high-throughput ML services

Fintech Co. · Software Engineer II    ~~(role eliminated)~~    *Jun-Sep 2026*  
SWE Co-op    *Sep-Dec 2025*

Added Sentry and Datadog observability to previously unmonitored services, including an MCP server, and prototyped AI-assisted incident triage

Cut table load time from 6s to under 1s with a decorator-based caching layer

Big Bank · Software Developer Intern (2 terms)    *2024*

SELECTED WORK

Multi-agent RL for cooperation, co-author, top-tier ML conference 2026 · [paper](#) · [code](#)

tinyglot: character-level Transformer, 10 languages · [live demo](#) · [write-up](#)

EDUCATION

BCS (Hons.), Co-op    **MOVED FROM THE TOP**    *2026*

~~Coursework: Distributed Systems, Networks, NLP, Neural Networks, RL~~

1 Title yourself as the role you want.  
Recruiters read the title and the file name before anything else. "New grad" says stage, not skill.

2 Lead with the strongest AI evidence.  
The LLM observability work was bullet 3 of job 2. For an ML infra role, it's the headline, with numbers.

3 Don't explain a gap on paper.  
If a Hiring Manager wants to know, they'll ask. Have a one-line answer ready.

4 Link everything.  
A paper or project nobody can click on counts for little.

5 Cut the coursework line.  
It takes space that shipped work should have. Education moves to the bottom.

THE AI ENGINEER PROFILE PLAYBOOK · RESUME

06 / 11

FIG 4.2 – MARCUS K. • DATA SCIENTIST → AI ENGINEERING LEAD

# Marcus K.

## Data Scientist

AI Engineering Lead · Production RAG & LLM systems · 8 yrs

Toronto, ON · marcus.k@email.com · github.com/marcusk

### SUMMARY

**Data Scientist and tech lead with 8 years of experience designing scalable AI & ML solutions...**

AI Engineering lead who ships LLM systems to production for enterprise clients and leads engineers. Has led teams of up to 6 and turned research-grade prototypes into evaluated, monitored services.

### EXPERIENCE

**Consultancy · AI Consultant (AI Engineering lead)** 2024 - now

- Built a production RAG architecture... developing ETL and chunking pipelines, embedding data into a vector store...**
- Led design and delivery of a production RAG system for a pharma client (Databricks, vector search); built a retrieval eval harness (recall@k, faithfulness) used to gate every release
- Shipped Text-to-SQL for actuaries on Azure, saving several hours per analyst each week; ran stakeholder workshops to scope it

**E-commerce Co. · Data Scientist** 2022 - 2024

- Consolidated four legacy LTV systems into one modular framework: \$5M/yr savings, guiding \$600M in spend

**Consultancy · Data Scientist (Tech Lead)** 2018 - 2022

- Led 6 engineers and BAs to ship an order distribution optimizer: \$3M/yr freight savings **ALSO IN SUMMARY**

### SKILLS

LLM systems: RAG, Text-to-SQL, LangGraph, evals and tracing

Engineering: Python, FastAPI, Docker, CI/CD, Spark, Airflow · Databricks, GCP, AWS, Azure

**Tableau, Looker, Plotly, Seaborn, D3.js, Hive**

**Relevant Courses: ML, NLP, Bayesian Statistics...**

1

The title is a filter.

To a recruiter, "Data Scientist" reads as "can't lead engineers or work with PMs". Name the role you're applying for.

2

Answer the biggest doubt first.

The proof of leading engineers was buried in 2018. It goes in the summary.

3

Show the engineering, not just the outcome.

Say how quality was measured. An eval harness is the most under-claimed skill there is. *Only claim it if you built it.*

4

Lead with the LLM stack.

BI-heavy tools pull the reader back toward "analyst".

5

No "Relevant Courses" at 8 years.

It undercuts the seniority.

FIG 4.3 — OMAR Y. • MSC RESEARCHER → AI ENGINEER

# Omar Y. Hassan

Was "Omar Y." here and "Omar H." on LinkedIn

~~MSc CS candidate — looking for research or engineering roles~~

AI Engineer • LLM security & evals • Eligible to work in Canada

Toronto, ON • omar.h@email.com • github.com/oy2001

## EXPERIENCE

**Gov. Research Lab • Agentic AI Security Engineer** *Oct 2024 – Mar 2025 • Remote*

- Built an agentic harness where LLMs plan and execute lateral-movement attacks; benchmarked 10 models × 5 temperatures across 3 autonomy modes

**AI Startup • Computer Vision Engineer (Deep Learning, VLMs)** *Aug 2024 – Aug 2025 • Cairo (on-site)*

- Shipped production detection and OCR models; built CLIP-style retrieval over 6M images

~~SOC Analyst Intern — Telecom Internship, 2022–23~~

~~Academic Experience: 4 research assistant roles...~~

## PROJECTS

- guardrail-bench: 11 LLM guardrail models vs 6 attack families • repo • write-up • thesis
- agent-redteam: open-source harness for testing tool-using agents • demo video
- rag-eval-kit: faithfulness and retrieval evals as CI checks • repo
- ~~Wireshark forensics — NYC taxi cleaning — University website backend — Fantasy football analysis~~

## EDUCATION

**MSc Computer Science (thesis: LLM guardrail security)** *2026*

1 One name everywhere.  
A different name on the resume than on LinkedIn means a recruiter can't connect the two.

2 Pick one: research or engineering.  
Research roles mostly need a PhD and are scarce in Canada. Pick the role the evidence supports, then specialize.

3 Explain overlapping jobs.  
Two countries at once looks like an error without "Remote". Keep official titles and add a descriptor, so it survives a reference check.

4 Cut off-target roles.  
It makes room for the specialty.

5 Make the thesis the projects.  
The AI security and evals work was one line under Education. Now it's three linked projects.

THE AI ENGINEER PROFILE PLAYBOOK • RESUME

08 / 11

# Weak, okay, strong

The same person's GitHub at three levels. Under each is what I conclude in the first 10 seconds.

- WEAK

The screenshot shows a GitHub profile for user 'oy2001'. The bio is '(no bio)'. The README.md file contains the text 'No profile README'. Under the 'Pinned' section, there are four repositories: 'cs3307-assignment3' (Java, 'No description'), 'transformers' (Python, 'Forked from huggingface/transformers'), 'dotfiles' (Shell, 'No description'), and 'project-final' (JavaScript, 'No description'). At the bottom, it shows '41 contributions in the last year' with a grid of green squares representing activity.

**What I conclude:** no evidence they build anything outside class.

~ OKAY

The screenshot shows a GitHub profile for user 'oy2001'. The bio is 'ML engineer · Toronto'. The README.md file contains the text 'No profile README'. Under the 'Pinned' section, there are four repositories: 'guardrail-bench' (Python, 'Benchmarking guardrails'), 'plate-ocr' (Python, 'License plate OCR'), 'rag-playground' (Python, 'No description'), and 'portfolio' (TypeScript, 'Personal site'). At the bottom, it shows '349 contributions in the last year' with a grid of green squares representing activity.

**What I conclude:** real projects and real activity, but the activity comes in bursts and mostly happened in the last month. It reads as "started building for the job hunt". It doesn't hurt, but it won't make you stand out.

+ STRONG

oy

oy2001

AI Engineer · LLM security & evals · Toronto

oy2001 / README.md

I build ways to test whether LLM systems are safe to ship.

Now: open-source red-team harness for tool-using agents

Latest write-up: [Your guardrail model has a blind spot](#)

Stack: PyTorch · vLLM · FastAPI · Docker · GitHub Actions

Pinned

guardrail-bench

11 guardrail models vs 6 attack families. Results table, write-up, 1-command repro.

Python ★ 48

agent-redteam

Red-team harness for tool-using agents. ▶ 2-min demo.

Python ★ 21

rag-eval-kit

Faithfulness and retrieval evals that run as CI checks.

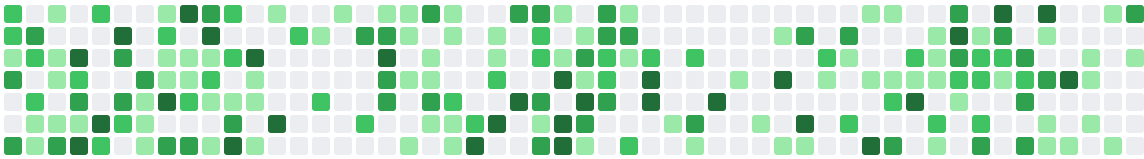
Python ★ 9

plate-ocr

License-plate OCR, from detection to recognition. Live demo.

Python ★ 6

612 contributions in the last year



**What I conclude:** a builder with a specialty. I know what to ask about before the interview starts.

## A project README, before and after

- BEFORE

```
# guardrail-bench

run: python main.py
```

**1 Nothing to judge.**

I can't tell what it does, whether it works, or why it matters, and I won't clone it to find out.

+ AFTER

```
# guardrail-bench
Which LLM guardrails actually stop attacks?
11 models × 6 attack families, one command.

Demo: 2-min video · Write-up · MSc thesis

## Results (attack success rate, lower is better)
| Guardrail | Jailbreak | Prompt inj. | Encoding |
|-----|-----|-----|-----|
| model-a | 4% | 31% | 58% |
| model-b | 9% | 12% | 22% |

## Run it
make bench MODEL=model-a

## How it works
attacks/ → harness/ → scorers/ → report.md
```

THE AI ENGINEER PROFILE PLAYBOOK · GITHUB

10 / 11

# Profiles worth studying

Exceptional, and not the bar for an interview. Look at the pattern: a clear one-liner, a few flagship projects, writing that explains them, steady output.

| PROFILE                                                         | WHAT TO NOTICE                                                                               |
|-----------------------------------------------------------------|----------------------------------------------------------------------------------------------|
| <a href="https://github.com/simonw">github.com/simonw</a>       | Ships small tools constantly (llm, datasette), and every one has a blog post explaining it.  |
| <a href="https://github.com/karpathy">github.com/karpathy</a>   | A few pinned repos (nanoGPT), each one minimal, readable and focused on teaching.            |
| <a href="https://github.com/rasbt">github.com/rasbt</a>         | A profile README that routes you to the work. LLMs-from-scratch is a model flagship project. |
| <a href="https://github.com/eugeneyan">github.com/eugeneyan</a> | Pairs repos with writing on applied ML. Proof that writing about your work counts.           |
| <a href="https://github.com/hamelsmu">github.com/hamelsmu</a>   | Practical writing on LLM evals, the most under-claimed skill on resumes.                     |

# Want these checks run on your profile?

This playbook is the sample. An audit applies every check to **your** LinkedIn, resume and GitHub, ranks the fixes by impact and picks one target role your evidence supports. I give the feedback; you make the changes.

| TIER                                              | INCLUDES                                                                                                                                                    | CAD     |
|---------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
| <b>Audit</b><br>WRITTEN · 72 H <b>RECOMMENDED</b> | Line-by-line audit of LinkedIn, resume and GitHub. Every fix ranked by impact. One target role and a top-3 project plan. 1 review of your revision.         | \$249   |
| <b>Rebuild</b><br>LIVE · ~2 WEEKS                 | Everything in the Audit, plus 2 × 45-min walkthroughs, GitHub pins and README guidance, a mock interview, and 2 reviews within 30 days.                     | \$749   |
| <b>Lead Rebuild</b><br>LIVE · SENIOR ROLES        | Everything in the Rebuild, plus a leadership story (leading engineers, working with PMs), system-design and behavioral mocks, and 2 reviews within 60 days. | \$1,250 |

Audit guarantee: fewer than 10 specific, actionable fixes and you get your money back. Book at [jateeq.github.io/ai-engineer-audit/](https://jateeq.github.io/ai-engineer-audit/) or message me on [LinkedIn](#).

Data: 17 resumes, 20 LinkedIn profiles and 13 GitHub accounts (63 showcased repos) from AI/ML job seekers, Sep-Oct 2026. GitHub figures come from public data via the GitHub API.

Every example is anonymized: names, employers, schools and numbers changed. Contribution graphs are illustrative.

Playbook v1.0 · October 2026 · Set in Inter, Source Serif 4 and JetBrains Mono.