

AI Engineer Profile Audit

Pass the screen · Stand out · Land the role

The AI Engineer Profile Playbook

Quick-fix checklists · real before/after · the mistakes I see most · profiles worth studying

By Jawad A., Lead AI Platform Engineer and Hiring Manager, who has interviewed 100s of AI/ML Engineers and reviewed thousands of profiles over almost a decade. This playbook is a **sample**: the fast, visible fixes I see most often, with real (anonymized) before/after. Tick each check that's true of your profile today.

The most common mistakes I see

From the resumes, LinkedIn profiles and GitHub accounts I broke down line by line.

Resume

Claims without proof	94%
No single target role	82%
Inconsistencies a Hiring Manager catches	82%
Best AI work buried	71%
Underselling in line one	47%
Evals built, never claimed	47%

LinkedIn

No GitHub link in Contact info	80%
Current role left blank or one line	60%
Headline isn't one AI/ML role	60%
No About section	30%
Top 3 skills unrelated to AI Engineering	25%

GitHub

Showcased repos with no CI	92%
Showcased repos with no tests	87%
Active less than half the year	77%
No profile README	69%

LinkedIn

LinkedIn: quick fixes

1 Your headline names one role, plus your edge

"AI Engineer · LLM evals & cloud infra", not five keywords separated by pipes.

2 The headline isn't a job request

"Seeking new grad roles" makes your job hunt the headline instead of what you can do. The green Open to Work frame is fine.

3 Featured shows something you built

Something a Hiring Manager can click and use: a live demo, a shipped app, real users. An empty Featured section is wasted space.

4 Every role has a description, especially the current one

60% of the profiles I sampled left their current role blank or one line. Recruiters search that text; 2-4 bullets with numbers, same as your resume.

5 Your name and titles match your resume

A different surname or a stale current role makes a recruiter wonder if it's the same person.

6 About opens with what you build

Not "passionate recent graduate". Three lines: what you build, the proof, what you want next.

7 Your top 3 skills are AI Engineering skills

They're pinned near the top, so they act as a second headline. "Figma, Product Strategy" or "Leadership, Entrepreneurship" pulls you away from the role.

8 No fancy Unicode fonts

"Bold" and "italic" look-alike characters pasted into About get stripped in exports and likely in search. They also break screen readers. Use plain text.

Before / after: a full profile

Meet Leila, a new grad with LLM agent research, eval work and production cloud experience. (Name and employers changed.)

BEFORE

Leila M.
~~Ex SWE @ Big Bank | ML & LLM Researcher | Seeking New Grad Roles~~ ¹
Toronto, Ontario, Canada · 500+ connections

About
~~Recent Computer Science graduate passionate about AI/ML, cloud and data engineering. Seeking new grad or junior opportunities where I can learn and grow...~~ ²

Featured
~~Nothing featured~~ ³

Experience

**Machine learning Research Assistant** ⁴
University · May 2025 – Aug 2026
Contributed to LLM-based recruitment systems focused on fairness...

AFTER

I build and evaluate LLM agents
Leila M.
AI Engineer · LLM agents & evals · AWS / Terraform in production
Toronto, Ontario, Canada · 500+ connections

About
I build LLM agents and the evals that tell you whether they work. Built an agentic recruitment assistant (LangChain, FastAPI) that 300+ job seekers have signed up for, with an eval pipeline that scores every release. Ran AWS infrastructure in production. Looking for an AI Engineer role on a team shipping LLM products.

Featured

**DEMO**
Recruiting agent
2-min demo + repo

**LIVE APP**
Try it yourself
Paste a job post, get a tailored resume

**SHIPPED**
App Store
300+ sign-ups, 4.7★

Experience

**Machine Learning Research Assistant**
University · May 2025 – Aug 2026 · Part-time
Built an agentic recruitment assistant and an eval pipeline over 100+ documents...

- 1. Leila's headline was labels and a job request.** "Ex", "Researcher" and "Seeking New Grad Roles" told recruiters where Leila had been and what she wanted, not what she can do. Now it names the role she's going after and what makes her different.
- 2. Her About section could have belonged to anyone.** "Passionate recent graduate" hid the agent Leila actually built. Now it opens with that system and backs it up in one line.
- 3. Her Featured section was empty.** It's the most visual spot on a profile, and Leila had real things to show. For an AI Engineer, that's what a Hiring Manager can try: a live app, a demo video, an App Store listing, user numbers. (Papers and press matter more for research roles.)
- 4. Her experience didn't say what she did.** Leila fixed the capitalization, marked the role part-time (it overlapped with her co-ops), and now opens with what she built.

Headline makeovers

Five real headlines (anonymized), each rewritten into one role plus an edge.

S

Sam A.
~~AI & ML Engineer | Machine Learning & Deep Learning |~~
~~Medical Imaging & Healthcare AI | Data Analyst |~~
~~Explainable AI~~

→

S

Sam A.
ML Engineer · Medical imaging & explainable deep learning

Five identities, including "Data Analyst", is the same as having none. Pick the one the experience supports.

M

Marcus K.
~~AI & ML @ Consultancy | M.S. Computer Science~~

→

M

Marcus K.
AI Engineering Lead · Production RAG & Text-to-SQL ·
Led teams of 6

The original names an employer and a degree. The fix names the role and the proof that answers "can he lead engineers?"

N

Nadia S.
~~Ex-Big 4 | Machine Learning Associate | MMAI '25~~

→

N

Nadia S.
AI Engineer · RAG & document AI · Shipped
containerized LLM APIs

"Ex-Big 4" leads with someone else's brand. The fix leads with what she has shipped.

O

Omar Y.
~~Agentic AI Security Engineer at Gov. Research Lab~~ (role
ended last year)

→

O

Omar Y.
AI Engineer · LLM security & evals · MSc thesis on
guardrails

A stale "current" role raises a question in the first second. The fix turns it into a specialty.

R

Rami A.
Senior Machine Learning Engineer @ Big Bank

→

R

Rami A.
Senior ML Engineer · LLM workflows & ML observability ·
Shipped a paid iOS app

The original is accurate but forgettable. The fix adds the two things that set him apart, including an app with paying users.

Resume

Resume: quick fixes

1 The title names the role you want

And so does the file name.

"Resume_SWE_New_Grad_V6.pdf" gets read too.

2 Your strongest AI work is in the top third

Not bullet 3 of job 2.

3 Every skill appears in a bullet

If "MLOps" lives only in the skills list, it reads as keyword stuffing.

4 Dates and titles survive a reference check

No unexplained overlaps. Keep official titles and add a descriptor instead of renaming.

5 Strong metrics state how they were validated

0.90 AUC on 86 samples needs "held-out subjects, 5-fold CV".

6 It reads cleanly as plain text

Paste your PDF into a text editor. If it's scrambled, applicant tracking systems scramble it too.

Before / after

Example: New grad → ML infra

BEFORE

Maya R.

~~Software Engineer - New Grad~~ ¹ · file:
~~Resume_SWE_New_Grad_V6.pdf~~

EDUCATION

BCS (Hons.), Co-op 2026

- ~~Coursework: Distributed Systems, Networks, NLP, Neural Networks, RL~~ ²

EXPERIENCE

Fintech Co. · Software Engineer II ~~(role eliminated)~~ ³ | SWE Co-op
Jun-Sep 2026, Sep-Dec 2025

- Built an end-to-end admin credential management feature that eliminated staff-assisted service-account revocation...
- Optimized table performance from 6s to <1s with a caching layer

Social Media Co. · Backend Engineering Intern, Ads ML
Jun-Sep 2025

- Built the first integration testing framework in Python for the ML platform...
- Architected a Prometheus-compatible metrics library to replace Graphite...
- Developed a Unified LLM Observability Tool, standardizing metric collection (latency, token count) across providers... ⁴

Big Bank · Software Development Intern ~~(2 x co-ops)~~
2024

RESEARCH & PROJECTS

- Promoting Multi-agent Co-operation, top-tier ML conference 2026 (no link) ⁵
- Character autocomplete Transformer, course project (no link)

AFTER

Maya R.

ML Infrastructure Engineer · LLM observability · Python, Go · github.com/mayar · mayar.dev

EXPERIENCE

Social Media Co. · Backend Engineering Intern, Ads ML Platform
Jun-Sep 2025

- Built a unified LLM observability library (latency, tokens, cost) across 3 model providers, adopted by 5 ML services for on-call debugging and cost tracking
- Designed the platform's first integration test framework (gRPC, permissioned access to real services), with Slack alerting on regressions
- Replaced Graphite with a Prometheus/OpenTelemetry metrics library for high-throughput ML services

Fintech Co. · Software Engineer II
Jun-Sep 2026
SWE Co-op Sep-Dec 2025

- Added Sentry and Datadog observability to previously unmonitored services, including an MCP server, and prototyped AI-assisted incident triage
- Cut table load time from 6s to under 1s with a decorator-based caching layer

Big Bank · Software Developer Intern (2 terms) 2024

SELECTED WORK

- Multi-agent RL for cooperation, co-author, top-tier ML conference 2026 · paper · code
- tinyglot: character-level Transformer, 10 languages · live demo · write-up

EDUCATION

BCS (Hons.), Co-op 2026

1. The "New grad" title and file name. Recruiters read the title and the file name before anything else. Title yourself as the role you want.

2. A coursework line. It takes space that shipped work should have.

3. Explaining a gap on paper. If a Hiring Manager wants to know, they'll ask. Have a one-line answer ready and keep the resume clean.

4. The strongest AI evidence was bullet 3 of job 2. The LLM observability work is the headline for an ML infra role, so it moves to the top.

5. No links. A paper or project nobody can click on counts for little. Every item gets a demo, code or a write-up.

Example: Data scientist → AI Engineering lead

BEFORE

Marcus K.

Data Scientist ¹

SUMMARY

~~Data Scientist and tech lead~~ with 8 years of experience designing scalable AI & ML solutions...

EXPERIENCE

Consultancy · AI Consultant 2024 - now

- Built a production RAG architecture... developing ETL and chunking pipelines, embedding data into a vector store, and optimizing retrieval to generate strategic briefs ²
- Implemented Text-to-SQL models enabling actuaries to query underwriting data...

E-commerce Co. · Data Scientist 2022 - 2024

- Led two LTV models consolidating four legacy systems, \$5M annual savings

Consultancy · Data Scientist 2018 - 2022

- Led a team of 6 software engineers and business analysts ³ to deliver an order distribution tool...

SKILLS

~~Tableau, Looker, Plotly, Seaborn, D3.js, Hive~~ ⁴,
LangChain, Spark, Airflow...

~~Relevant Courses: ML, NLP, Bayesian Statistics...~~ ⁵

AFTER

Marcus K.

AI Engineering Lead · Production RAG & LLM systems
· 8 yrs

SUMMARY

AI Engineering lead who ships LLM systems to production for enterprise clients and leads engineers. Has led teams of up to 6 and turned research-grade prototypes into evaluated, monitored services.

EXPERIENCE

Consultancy · AI Consultant (AI Engineering lead)

2024 - now

- Led design and delivery of a production RAG system for a pharma client (Databricks, vector search); built a retrieval eval harness (recall@k, faithfulness) used to gate every release
- Shipped Text-to-SQL for actuaries on Azure, saving several hours per analyst each week; ran stakeholder workshops to scope it

E-commerce Co. · Data Scientist 2022 - 2024

- Consolidated four legacy LTV systems into one modular framework: \$5M/yr savings, guiding \$600M in spend

Consultancy · Data Scientist (Tech Lead) 2018 - 2022

- Led 6 engineers and BAs to ship an order distribution optimizer: \$3M/yr freight savings

SKILLS

LLM systems: RAG, Text-to-SQL, LangGraph, evals and tracing · **Eng:** Python, FastAPI, Docker, CI/CD, Spark, Airflow · Databricks, GCP, AWS, Azure

1. **"Data Scientist" in the title.** To a recruiter, that reads as "can't lead engineers or work with PMs". The title is a filter, so name the role you're applying for.
2. **Business outcomes only, no engineering signal.** Add how quality was measured. The eval harness is the most under-claimed skill there is. *Only claim it if you built it.*
3. **The proof of leading engineers was buried in 2018.** It answers the recruiter's biggest doubt, so it goes in the summary.
4. **BI-heavy tools.** They pull the reader back toward "analyst". Lead with the LLM and engineering stack.
5. **A "Relevant Courses" line on an 8-year resume.** It undercuts the seniority.

Example: MSc researcher → AI Engineer

BEFORE

Omar Y. ~~(LinkedIn: Omar H.)~~ ¹

MSc CS candidate · ~~looking for research or engineering roles~~ ²

INDUSTRY EXPERIENCE

Gov. Research Lab · Agentic AI Security Engineer

Oct 2024 - Mar 2025 · Ontario

AI Startup · Computer Vision Engineer

Aug 2024 - Aug 2025 · Cairo (on-site) ³

~~SOC Analyst Intern - Telecom Internship~~ ⁴ 2022-23

ACADEMIC EXPERIENCE

- 4 research assistant roles...

PROJECTS

- ~~Wireshark forensics - NYC taxi cleaning - University website backend - Fantasy football analysis~~ ⁵

Thesis: LLM guardrail security (one line under Education)

AFTER

Omar Y. Hassan

AI Engineer · LLM security & evals · Eligible to work in Canada

EXPERIENCE

Gov. Research Lab · Agentic AI Security Engineer

Oct 2024 - Mar 2025 · Remote

- Built an agentic harness where LLMs plan and execute lateral-movement attacks; benchmarked 10 models × 5 temperatures across 3 autonomy modes

AI Startup · Computer Vision Engineer (Deep Learning, VLMs)

Aug 2024 - Aug 2025

- Shipped production detection and OCR models; built CLIP-style retrieval over 6M images

PROJECTS

- [guardrail-bench](#): 11 LLM guardrail models vs 6 attack families · [repo](#) · [write-up](#) · [thesis](#)
- [agent-redteam](#): open-source harness for testing tool-using agents · [demo video](#)
- [rag-eval-kit](#): faithfulness and retrieval evals as CI checks · [repo](#)

EDUCATION

MSc Computer Science (thesis: LLM guardrail security)

2026

- 1. A different name on the resume than on LinkedIn.** A recruiter can't connect the two. Use one consistent display name everywhere.
- 2. "Research or engineering."** Research roles mostly need a PhD and are scarce in Canada. Pick the role the evidence supports, then specialize.
- 3. Overlapping jobs in two countries.** Without "Remote", it looks like an error. Keep the official title and add a descriptor rather than renaming it, so it survives a reference check.
- 4. Off-target internships.** Cutting them makes room on the page.
- 5. Generic projects while the thesis sat in one line.** The AI security and evals work is the specialty. Make it the three featured projects, each with a link.



GitHub: quick fixes

- 1 A one-line bio and a profile README**
What you build, in one sentence. That's the first thing I read.
- 2 2-4 pinned projects you're proud of**
Not assignments, forks or dotfiles. Each one needs a description that says what it does and the result.

- 3 READMEs lead with the result**
The question, the result and a demo. Then how to run it.
- 4 Activity shows steady building**
Most weeks, all year. A burst right before the job hunt reads as exactly that.

Weak → okay → strong

Contribution graphs are illustrative. The "okay" one is modeled on a real profile.

WEAK

oy

oy2001
(no bio)

README.md
No profile README

Pinned

cs3307-assignment3
No description
● Java

transformers
Forked from huggingface/transformers
● Python

dotfiles
No description
● Shell

project-final
No description
● JavaScript

41 contributions in the last year

What I conclude: no evidence they build anything outside class.

OKAY

oy

oy2001
ML engineer · Toronto

README.md
No profile README

Pinned

guardrail-bench
Benchmarking guardrails
● Python

plate-ocr
License plate OCR
● Python

rag-playground
No description
● Python

portfolio
Personal site
● TypeScript

349 contributions in the last year

What I conclude: real projects and real activity, but the activity comes in bursts and mostly happened in the last month. It reads as "started building for the job hunt". It doesn't hurt, but it won't make you stand out.

STRONG

oy

oy2001
AI Engineer · LLM security & evals · Toronto

oy2001 / README.md
I build ways to test whether LLM systems are safe to ship.
Now: open-source red-team harness for tool-using agents
Latest write-up: [Your guardrail model has a blind spot](#)
Stack: PyTorch · vLLM · FastAPI · Docker · GitHub Actions

Pinned

guardrail-bench
11 guardrail models vs 6 attack families. Results table, write-up, 1-command repro.
● Python ★ 48

agent-redteam
Red-team harness for tool-using agents. ▶ 2-min demo.
● Python ★ 21

rag-eval-kit
Faithfulness and retrieval evals that run as CI checks.
● Python ★ 9

plate-ocr
License-plate OCR, from detection to recognition. Live demo.
● Python ★ 6

612 contributions in the last year

What I conclude: a builder with a specialty. I know what to ask about before the interview starts.

A project README, before and after

BEFORE

```
# guardrail-bench  
  
run: python main.py
```

1. **Nothing to judge.** I can't tell what it does, whether it works, or why it matters, and I won't clone it to find out.

AFTER

```
# guardrail-bench  
Which LLM guardrails actually stop attacks?  
11 models × 6 attack families, one command.  
  
Demo: 2-min video · Write-up · MSc thesis  
  
## Results (attack success rate, lower is better)  
| Guardrail | Jailbreak | Prompt inj. | Encoding |  
|-----|-----|-----|-----|  
| model-a | 4% | 31% | 58% |  
| model-b | 9% | 12% | 22% |  
  
## Run it  
make bench MODEL=model-a  
  
## How it works  
attacks/ → harness/ → scorers/ → report.md
```

Profiles worth studying

These are exceptional, and they're not the bar for getting an interview. Look at the pattern: a clear one-liner, a few flagship projects, writing that explains them, and steady output.



simonw

Ships small tools constantly (llm, datasette), and every one has a blog post explaining it.



karpathy

A few pinned repos (nanoGPT), each one minimal, readable and focused on teaching.



rasbt

A profile README that routes you to the work. LLMs-from-scratch is a model flagship project.



eugeneyan

Pairs repos with writing on applied ML. Proof that writing about your work counts.



hamelsmu

Known for practical writing on LLM evals, the most under-claimed skill on resumes.

Want me to run these checks on your profile?

Book an audit: cal.com/jawad-ateeq-hgbvwe/audit-ai-engineer-profile · More at jateeq.github.io/ai-engineer-audit/